

## Project 2 Report (Thread-safe Malloc)

### 1. Implementation of thread-safe malloc

In order to prevent race conditions in the implementation of malloc and free, which come from the execution of multi-threaded code due to different timing conditions, I use two strategies to solve this problem, including locking version and non-locking version.

#### 1.1 Locking version

I created a LinkedList to implement malloc and free, and I introduced two global variables to track the head and the tail of the free LinkedList, which are `head_free_metadata_lock` and `tail_free_metadata_lock`. In multiple threads, the global variables should be locked in a thread in order to keep them from being affected by another thread, so I add locks to these variables. In particular, I initialize the mutex called `lock` in

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

Then I add

```
pthread_mutex_lock(&lock);  
pthread_mutex_unlock(&lock);
```

for `bf_malloc(size, &head_free_metadata_lock, &tail_free_metadata_lock, 1)` and `bf_free(ptr, &head_free_metadata_lock, &tail_free_metadata_lock)` since they call these two variables. Between `pthread_mutex_lock` and `pthread_mutex_unlock` functions is the place called the critical section, which allows access to a shared resource by no more than one thread at the same time. I put `bf_malloc` and `bf_free` in the critical section, which can allow concurrency and also prevent race conditions. The last parameter in `bf_malloc` is another variable called `isLock` to help identify whether I should lock the function `sbrk` or not. For the lock version, `isLock` is equal to 1, the variables of the entire function `bf_malloc` are locked when they are called, so they don't need another lock for `sbrk` to ensure the thread-safe feature. When we call the function `increaseHeapMemory(size_t size, int isLock)`, we will call `sbrk` function, but we don't need to add a lock to it.

#### 1.2 Non-locking version

In the non-locking version of thread-safe model, I use Thread-Local Storage(TLS) to make sure the resources in one thread will not be affected by another thread, since TLS allocates separate memory for variables in different threads. So I add the keyword `__thread` in front of the global variables `head_free_metadata_nolock` and `tail_free_metadata_nolock` to implement TLS, which allows concurrency and also prevent race conditions because they don't have the shared global variables. So we don't add `pthread_mutex_lock` and `pthread_mutex_unlock` to `bf_malloc(size,`

`&head_free_metadata_nolock, &tail_free_metadata_nolock, 0)` and `bf_free(ptr, &head_free_metadata_nolock, &tail_free_metadata_nolock)`. For the `bf_malloc` function, the variable `isLock` is equal to 0, which means we should use a lock before calling `sbrk` and release a lock after calling `sbrk` because it's not thread-safe. So when we call the function `increaseHeapMemory(size_t size, int isLock)`, we will call `sbrk` function, it's where we should apply a lock to the `sbrk` function.

## 2. Performance result and comparison of locking vs. non-locking version

The results below show the execution time and data segment size for the locking version and non-locking version. As we can see, the execution time and data segment size of the locking version are slightly larger than that of the non-locking version. Locking version applies mutex for the variables in the `bf_malloc` function, where non-locking version only applies mutex for the function `sbrk`. The locking version has more resources locked in multi-threaded code, which will take longer time for another thread to execute. The non-locking version is more efficient for memory allocation, since more resources are allocated and freed at the same time, which can utilize the space where is just freed or allocated before. The non-locking version is slightly better compared with the locking version given these two measurements.

### Locking version:

```
ly161@vcm-24863:~/ECE650/hw2/thread_tests$ ./thread_test_measurement
No overlapping allocated regions found!
Test passed
Execution Time = 0.132278 seconds
Data Segment Size = 44484064 bytes
```

### Non-locking version:

```
ly161@vcm-24863:~/ECE650/hw2/thread_tests$ ./thread_test_measurement
No overlapping allocated regions found!
Test passed
Execution Time = 0.123047 seconds
Data Segment Size = 43320096 bytes
```